

MOHAMMED ALKHAZRAJI
25250483
CHATGPT USED FOR ALL PARTS



PART 1


DOMAIN : MARKETING

Task Description: The task is to generate a marketing scheme to promote a new product using Instagram.

VERSION A :

THE PROMPT : Give me a marketing scheme for Instagram.





VERSION B :

THE PROMPT :

Act as an Instagram marketing expert

I need you to write a comprehensive marketing campaign plan for a small fast food truck business called DR,Burger

Target audience : students and young professionals who wants to grab fast food to ease and fasten their life

MAIN GOALS :

build brand attention and focus on virality






Highlight the fresh and locally sourced ingredients

THE PLAN MUST INCLUDE :

Content themes : what topics to post about e.g ' menu items , customer reviews '

Hashtag strategy : provide 10 to 15 specified hashtags

Techniques used : role prompting , structured output , detailed context



Comparison Analysis :

The difference in quality , usefulness and specificity is huge , the output from version a was generic , unusable and offering a list of known things to do , version b on the other hand was actionable and highly specific.

The superiority of version b is clear in three concrete improvements

- 1 audience specified content
- 2 structured format
- 3 tailored brand strategy

Version a failed because it provided the AI with no guidance

Version B was better because it used several prompt engineering concepts :

Providing Rich Context: the specified prompt provided a name a strategy a location and a theme unlike the vague prompt which was generic and unuseful

Defining Audience & Tone : version b offered a key concept which is audience defining to target the audience wanted for the food truck project which gave the prompt a more clear answer because of the context provided

Role Prompting : version B used role prompting (act as) this technique frames the AI mindset to give the output based on the related subject

Enforcing a Structural Contract : the specified prompt offered demanded a constant schedule and specific content like themes and hashtag strategies by this the AI is forced to provide usable and comprehensive document\

PART 2

For this assignment, I have selected Task 1: Design a customer onboarding email sequence. The framework I have chosen to do this is the CO-STAR framework.

This framework is best for this particular job because it allows for a direct and explicit statement of the creative brief. For example, to create effective emails, one needs to maintain a strong, consistent brand voice. components of framework example style and tone components within CO-STAR directly deal with the former and allow me to teach the AI how to sound, not just what to say. The another component is audience component will be crucial for an onboarding sequence since the AI needs to understand who it's speaking to a new customer, thus this framework brings in the much required structure that will lead one to a high quality and brand aligned creative output.

The prompt :

CONTEXT : You are the new Customer Experience Manager for "MGS," an online subscription service which sends curated indoor plants and pots to beginners. A new customer has just made their first purchase-the "Beginner's Box"-and signed up for an account.

OBJECTIVE : The design of a 3-email automated customer onboarding sequence to welcome the new customer, create excitement while their box is shipping, and provide some initial value to reduce anxiety and make them confident in this new plant-parenting hobby.

STYLE : Clear, helpful, and concise. Use short paragraphs, simple language, and bullet points where helpful. Each email must have one single, clear call-to-action.

TONE : Warm, encouraging, enthusiastic. The tone should be like a knowledgeable friend who is excited to help them succeed, not a cold corporation. Use "you" and "we" to create a personal connection.

AUDIENCE : This is a brand new customer. Most likely a total beginner to indoor plants. They are excited about their purchase, but also nervous they will accidentally kill their new plants. They want reassurance and easy, unintimidating guidance.

Please write the complete text for all 3 emails in the chain. Every email should include:

A Compelling Subject Line.

A custom greeting




The whole body text.

One clear call-to-action text per button

Email 1 - Trigger: Immediately after the purchase, focus on welcoming them and confirm the order.

Email 2 - (Trigger: 3 days post-purchase): Emphasize building excitement and provide one simple "get ready" tip.

Email 3 - Trigger: 7 days after purchase or 1 day after estimated delivery Basically, this email is all about "unboxing" and finding the right spot for their new plant.



PART 2 D :

Here are two clear weaknesses from the AI output

Very Long Subject Lines: The subject lines in all three emails are decidedly too long-for instance, Your plants are almost here, get ready to welcome them. On a mobile phone, they will be truncated, and the message will be lost.

Generic Brand Voice: It's friendly, but it doesn't have much character. It sounds like any nice company. It doesn't have a distinctive, memorable personality. It is helpful but not charismatic.

This is a great start, but I need to refine it. My two main changes are the following:


The refined prompt :

Refine the Subject Lines: All three subject lines are too long and will be cut off on mobile.

Please rewrite them to be shorter, punchier, and under 7 words each use more personality:

The brand voice is friendly but a little general let's make it more distinct. I want the "MGS" voice to sound like a plant obsessed best friend who is super excited for you. Let's call this persona the Plant Coach

Please recreate the complete 3 email series with these two changes.



What Changed:

The subject lines went from what was probably something long and boring, like, "Your MGS Beginner's Box Order is Almost Ready," to short, exciting, and emoji-driven phrases:

Why it's better: Nobody reads long subject lines, especially on a phone. These new versions are built for mobile. They create excitement and curiosity, which massively increases the chance of the email actually being opened.

The Brand Voice

What was different: It went from the standard, polite e-commerce bot voice into a character. This new "Plant Coach" character is upbeat, supportive, and a little quirky.

Why it's Better: A strong brand voice builds a relationship, not just a sale.

It stirs up excitement , community and trust

Framework Effectiveness Reflection :

It was very effective to work within the refining framework. It gave a clear, logical structure that moved easily from problem to solution better than a general make it better kind of brief, the format forced me to first identify the problems-specifically, long subjects and generic voice and then to provide clear, actionable solutions

For this specific task of refining an existing, weak output, I wouldn't use a different framework; it's perfect for the job. If starting from scratch, I may have used a different framework, but to iteratively improve this is the most direct route.

What I learned is that the framework's value really has to do with context.refining is a surgical tool for improvement

PART 3 :

At the start of this course, my idea of prompt engineering was very basic: basically, asking the AI a question-albeit one which in most cases yielded generic and useless answers like the first "vague" prompt that I tested in Part 1.

Now, I understand that prompt engineering means the ability to provide exact, formal, and highly structured instructions; it is not asking a question, but designing the answer. My

biggest "aha moment" came with Part 2 of this midterm, which applied the lessons from Week 5, "Controlling Style and Tone". I realized that I could control not just what it said, but how it said it.

I was engineering a personality, not just a text. The

most difficult thing so far has been Hallucinations (Week 2), the idea that the AI is confidently wrong. I overcame this by learning to ground the AI, meaning I provide the facts myself

Practical Application

A key real-world scenario in my web development career is generating technical documentation. When I finish a complex function or API endpoint, I need to write clear docstrings and README explanations—a time-consuming process. I could use prompt engineering to generate accurate first drafts of this documentation. I would primarily use the Context-Role-Task framework,I would give the code snippet as the Context, the Role of an "expert technical writer specializing in

Python docstrings," and a specific Task: "Generate a Google-style docstring for this function, including all parameters, return types, and potential exceptions."

The main risk, though, is a technical inaccuracy: AI may get the logic of the code wrong or miss an edge case. To address this, I would treat the output as a draft and not a final product. I would always manually review the generated docs for accuracy, using my knowledge at the developer level as the final human-in-the-loop check.

Limitations & Ethics

Two of the main limitations when prompting include hallucinations and a lack of true context. An LLM might confidently make up technically incorrect details of a function that sound real. It doesn't understand the business logic behind my code, just its syntactical patterns.

The major ethical concern that would be considered in my developer scenario is that of data privacy and intellectual property. Feeding my company's proprietary source code into a public, third-party AI model constitutes a major data leak and a breach of confidentiality.

I would enforce the following two policies: data sanitization and human-in-the-loop. First, I would make sure that all code snippets are generalized and anonymized, removing all sensitive IP or API keys. Second, human-in-the-loop: I always serve as the final validator who shall proofread every single line of generated documentation for technical accuracy before committing.

Goals for the Second Half

Aspects to Investigate Further

Going forward, I want to explore advanced reasoning and problem-solving patterns. Thus, I've only used prompts for simple generation: things like documentation. I'm far more interested in using AI as a collaborative partner in solving complex tasks. I want to learn about specific frameworks like CoT or ReAct that will help guide the AI through debugging complex logical errors, architecting new software modules, or optimizing inefficient code. I want to see how to make the AI "show its work" so I can validate its logic, not just its final answer.

Skill Gap to Close The big skill gap that I want to close is how to move from single-shot prompting to iterative, multiturn refinement. Presently, I can write a good prompt to get a decent first answer. But the real skill in development is handling complex, "wicked"

problems that have no single right answer. My goal is to become proficient in using a conversation with the AI to solve these problems-knowing how to provide feedback on its poor suggestions, how to ask it to "think" about a problem from a different angle, and how to collaboratively guide it toward a production-ready solution.