

Answer Key

Final Examination

Whole course · 40 questions × 2.5 points = 100 points

Answer string: B B B B B B D B B B C C C B A C C B B B B B C B C B C B D B B
D B B B B C B B B

1. **B** A probability distribution over the vocabulary; one token selected; repeat.
2. **B** Plausibility is optimised and nothing checks truth. It is the design working normally where plausible and true diverge.
3. **B** Thin coverage plus a specific detail that must be recalled rather than inferred. Citation format is highly patterned and easy to imitate.
4. **B** The likely continuation accepts the framing. The system completes your question rather than auditing it.
5. **B** Stacked suffixes produce rare assembled forms, so they fragment into several pieces.
6. **B** Tokenisation. Letters were never perceived, so it is being asked about something outside its representation.
7. **D** Training data is not in the context. The context is what is visible at answering time.
8. **B** Deliberate randomness in token selection; identical prompts can diverge.
9. **B** Consistency measures pattern strength. Widely repeated errors reproduce very consistently.
10. **B** They correlate because the training text was mostly written by people trying to be accurate. The course concerns the cases where they come apart.
11. **C** Instructions change tone, never available information. Requesting authority yields authority.
12. **C** Wrong shape means Format — the cheapest lever and the most neglected.
13. **C** Specific and checkable. The others are complaints rather than constraints.
14. **B** Displacement: nothing false, nothing applicable.
15. **A** Fabrication: an invented specific inside an otherwise correct answer.
16. **C** Omission. There is no wrong sentence to point at, so only prior knowledge catches it.
17. **C** Answered from memory. Absence of a quotation is the signature.
18. **B** Over-reading: the claim is stronger than the sentence supports.
19. **B** Retrieved the wrong passage — a genuine quotation from an unrelated place.
20. **B** Silent failure: wrong input, confident output, no signal. The characteristic multi-step failure.
21. **B** All nine are invisible in the output. This is why the response is procedural rather than perceptual.

22. **B** Consistent error resembles reliability. Random failure would be noticed.
23. **C** The quote test. A poisoned change is not supported by any real passage, so verification doubles as a security control.
24. **B** The unconscious checking of every intermediate output — free quality control that is deleted without being noticed.
25. **C** $0.95^{10} \approx 0.60$.
26. **B** Errors are inherited by later steps rather than averaged away.
27. **C** Narrow tools is a wall; the other three are requests that depend on good behaviour.
28. **B** Before the irreversible action. Identify the irreversible step first.
29. **D** Difficulty is not a criterion — and intellectually hard tasks are poor candidates. The best candidates are annoying rather than difficult.
30. **B** “A student misses a resit” rules out designs that “users could get wrong information” permits.
31. **B** Capability is free; constraint is the engineering.
32. **D** A Domain Pack does nothing for accuracy.
33. **B** Professional formatting increases persuasiveness without touching correctness, so it can increase over-trust.
34. **B** Each structure makes a particular error visible — SOAP separates report from measurement; IRAC states the rule before applying it.
35. **B** One undifferentiated token stream with no reliable label distinguishing instruction from data.
36. **B** The user never sees it and had no opportunity to decline. Direct injection is a user attacking their own tool; indirect injection has the user as victim.
37. **C** Causing an action requires tools. This is why tool scope is a security decision.
38. **B** Removing names is not anonymisation; a rare combination of ordinary details identifies people in small populations.
39. **B** The question is what is measured and what is certified — not whether AI was involved or whether use is detectable.
40. **B** The judgment that catches errors is built by doing the work. Without it you cannot evaluate what you are given.

Marking. 2.5 points per question, 100 points total, 60% of the course grade. No negative marking.