

Answer Key

Makeup Examination

Whole course · 40 questions × 2.5 points = 100 points

1. **B** Next-token prediction, repeated. Nothing is retrieved and no rules are applied.
2. **B** The same operation runs regardless of pattern strength, and confidence is not part of what is produced.
3. **B** It answers from the earlier version or something adjacent, without flagging the cut-off.
4. **D** Accepting a false premise follows from predicting a likely continuation, not from tokenisation. The other three are direct tokenisation consequences.
5. **C** Rare assembled forms fragment; common short words do not.
6. **B** Silent truncation with a confident answer. An explicit error would be the preferable case.
7. **B** Safety fine-tuning shifts probabilities. It produces a tendency, not a gate — which is why framing can defeat it.
8. **B** Deliberate randomness in token selection.
9. **C** Wrong reader or wrong country is a Context failure.
10. **B** A quotation can be searched for in the source. It does not prevent hallucination — it makes it detectable.
11. **B** It describes a feeling, not a text; many outputs satisfy it, so results vary.
12. **C** Field is subject vocabulary, tenor is relationship, mode is channel.
13. **B** The fallback principle: make the honest answer probable rather than appealing to honesty.
14. **C** Omission — true throughout, missing what mattered.
15. **A** Displacement — accurate elsewhere, inapplicable here.
16. **B** Retrieved the wrong passage: a genuine quotation from an unrelated place.
17. **A** Over-reading: an exception reported as a general rule.
18. **B** Silent failure.
19. **B** All are invisible in the output, which is why the response must be procedural.
20. **C** Verification requires something outside the system. The other three stay inside it.
21. **B** Correct only for the cases tried — untested paths were generated but never run.
22. **B** Consistency resembles reliability. Intermittent failure attracts attention; steady error does not.
23. **B** Inspection of every intermediate output — free quality control, deleted without notice.
24. **C** $0.95^5 \approx 0.77$.

- 25. **B** Completion is a judgment about matching an intention that is not present in the context, so it is predicted like anything else.
- 26. **C** Narrow tools is the only control that does not depend on the system behaving well.
- 27. **C** Sending is irreversible and no human precedes it.
- 28. **B** Repetitive, document-backed, recoverable. Difficulty is not a virtue here.
- 29. **B** It commits to nothing and changes no design decision. A named person does.
- 30. **B** Capability is free; constraint is the engineering.
- 31. **D** A Domain Pack does nothing for truth.
- 32. **B** It increases persuasiveness without improving correctness, so over-trust rises.
- 33. **B** Stating the rule before applying it makes a missing or wrong rule visible.
- 34. **B** Instructions and data share one undifferentiated channel, with no reliable trust label.
- 35. **B** Indirect injection — hidden in content the user never saw.
- 36. **C** Causing an action requires connected tools; the others need only text.
- 37. **B** A shared link carries the whole history. Screenshots crop; links do not.
- 38. **B** Rare combinations of ordinary details identify people in small populations.
- 39. **B** It breaks the measurement the assessment was designed to make. The problem is the broken measurement, not the tool.
- 40. **B** The judgment that lets you evaluate generated work is built by producing work yourself.

Marking. 2.5 points per question, 100 points total. No negative marking.