

Prompt Engineering and Large Language Models

Quiz 4

Weeks 7–8: Grounding, and Verification

Name: _____ Student No: _____ Date: _____

Instructions. This paper contains 10 multiple-choice questions. Select the **single best** answer for each. Each question is worth equal marks. Time allowed: 10 minutes. No notes or electronic devices.

1. “Grounding” a system in your own sources means:
 - A) Retraining the model on your documents
 - B) Placing relevant passages from your documents into the context when it answers
 - C) Restricting the model to a single topic
 - D) Reducing the model’s temperature
2. The abbreviation RAG stands for retrieval-augmented generation. In practice it means:
 - A) A method for training smaller models
 - B) Retrieve relevant text, add it to the context, then generate
 - C) A technique for reducing token costs
 - D) A way of checking output against the internet
3. Which problem does grounding **not** solve?
 - A) The model has never seen your institution’s internal policy
 - B) The guideline was updated after the knowledge cut-off
 - C) Your source document is itself out of date
 - D) Turkish-language professional material is thinly represented
4. Why should the refusal instruction specify a fixed string such as “NOT IN SOURCES”?
 - A) It reduces the number of tokens used
 - B) It is easy to spot when skim-reading, unlike a hedge buried in a paragraph
 - C) It prevents the model from answering at all
 - D) It is required by most document tools
5. An assistant answers “The notice period is 30 days” with no quotation from your documents. Which grounded failure mode should you suspect first?
 - A) Over-read the passage
 - B) Retrieved the wrong passage
 - C) Answered from memory
 - D) Silent failure
6. Your source states: “Either party *may* request an extension.” The assistant reports: “Either party *must* request an extension.” This is:
 - A) Fabrication

- B) Over-reading the passage
 - C) Answered from memory
 - D) Displacement
7. In a ten-question test set, what is the recommended composition?
- A) Ten questions your sources answer
 - B) Seven your sources answer and three they cannot
 - C) Five easy and five hard
 - D) Ten questions chosen at random
8. Why do the unanswerable questions matter more than the answerable ones?
- A) They are harder to write
 - B) They test behaviour in the situation you cannot detect during normal use
 - C) They take less time to grade
 - D) They are more likely to appear in the exam
9. Why must you write the expected answer **before** running each test question?
- A) It saves time
 - B) Otherwise you will unconsciously accept whatever output you receive
 - C) The tool requires it
 - D) It improves retrieval accuracy
10. Which of these is genuine verification?
- A) Asking the same question again and getting a similar answer
 - B) Asking the model whether it is confident
 - C) Searching your source document for the quoted sentence
 - D) Regenerating the answer three times and taking the majority