

Prompt Engineering and Large Language Models

Quiz 5

Weeks 9–12: Agents, Design, Domain Adaptation, and Security

Name: _____ Student No: _____ Date: _____

Instructions. This paper contains 10 multiple-choice questions. Select the **single best** answer for each. Each question is worth equal marks. Time allowed: 10 minutes. No notes or electronic devices.

1. What is the most important thing you lose when you replace a chatbot with an agent?
 - A) Speed
 - B) The ability to change the prompt
 - C) The checking you were performing on every intermediate output without noticing
 - D) Access to the context window
2. If each step of a workflow is 95% reliable, roughly what is the end-to-end reliability over ten steps?
 - A) 95%
 - B) 85%
 - C) 60%
 - D) 5%
3. Why is compounding error in an agent worse than independent random error?
 - A) Agents run more slowly
 - B) Each step consumes the previous step's output, so early errors are inherited
 - C) Agents use more tokens
 - D) Errors are reported only at the end
4. A three-step workflow is given faulty input at step one and produces a complete, confident, well-formatted result with no warning. This is:
 - A) Wrong stopping point
 - B) Silent failure
 - C) Context overflow
 - D) Indirect injection
5. Which control does **not** depend on the system behaving correctly?
 - A) Instructing it to be careful
 - B) Asking it to check its own work
 - C) Giving it only the narrowest tools needed for the job
 - D) Requesting quotations
6. In the Week 10 design worksheet, where should the human sit?
 - A) At the end, reviewing the final output

- B) Before the irreversible action
 - C) At the beginning, writing the prompt
 - D) Wherever there is time
7. Why is naming a specific harmed person better than naming a category of user?
- A) It is required by the university
 - B) Specificity changes the design; abstraction does not
 - C) It makes the document longer
 - D) Categories are difficult to define
8. A Domain Pack contains a glossary, golden examples, and output conventions. What does it **not** improve?
- A) Local vocabulary
 - B) Recognisable structure
 - C) Factual accuracy
 - D) Professional restraint
9. Why is indirect prompt injection more dangerous to you than direct injection?
- A) It is technically more sophisticated
 - B) The instruction is hidden in content you never see, so there is no point at which you could refuse
 - C) It only affects agents
 - D) It cannot be detected by any means
10. What is the root cause of prompt injection?
- A) Insufficient safety training
 - B) The model receives one stream of tokens and cannot reliably distinguish instructions from data
 - C) Context windows are too small
 - D) Users write careless prompts